

Lab 2. Gradient descent and Linear regression

Intro to Machine Learning
Fall 2018, Innopolis University

Lecture recap

- What is linear regression?
- Why study linear regression?
- What can we use it for?
- How to perform linear regression?
- How to estimate its performance?
- What are its extensions?
- What are dummy variables?

Questions about the lecture

Was the material already familiar to you?

What new things have you learned?

What was hard to understand?

What is p-value?


What is R^2 metric?

What is F-metric?


Today's plan

1. Gradient descent
 - a. Introduction, purpose and derivations
 - b. Types of GD
2. Linear regression
 - a. How it works
 - b. Implementation
3. Introduction to Scikit-learn
 - a. Linear model Implementation
4. Homework explanation

Gradient descent



Linear case



$$y = \beta_1 x + \beta_0$$

$$J = MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

$$J = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - (\beta_1 x + \beta_0))^2$$


Gradient descent step

$$\frac{\partial J}{\partial \beta_1} = -\frac{1}{n} \sum_{i=1}^n 2x_i(\hat{y}_i - (\beta_1 x_i + \beta_0))$$



$$\frac{\partial J}{\partial \beta_0} = -\frac{1}{n} \sum_{i=1}^n 2(\hat{y}_i - (\beta_1 x_i + \beta_0))$$

$$\beta_1 = \beta_1 - \alpha \frac{\partial J}{\partial \beta_1}$$

$$\beta_0 = \beta_0 - \alpha \frac{\partial J}{\partial \beta_0}$$




GD visualization



Different types of GD


1. Batch GD
2. Mini-batch GD
3. Stochastic GD



Try to guess
which is which?

Different types of GD

1. Batch GD
2. Mini-batch GD
3. Stochastic GD



— Batch gradient descent
— Mini-batch gradient Descent
— Stochastic gradient descent

Linear regression

$$RSS = e_1^2 + e_2^2 + \dots + e_n^2$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{y}_i = \hat{\beta}_1 x_i + \hat{\beta}_0$$

$$e_i = y_i - \hat{y}_i$$

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i, \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

Let's find β_0

$$RSS = (y_1 - \hat{\beta}_1 x_1 - \hat{\beta}_0)^2 + (y_2 - \hat{\beta}_1 x_2 - \hat{\beta}_0)^2 + \dots + (y_n - \hat{\beta}_1 x_n - \hat{\beta}_0)^2$$

$$\frac{\partial}{\partial \hat{\beta}_0} RSS =$$

$$\frac{\partial}{\partial \hat{\beta}_0} [(y_1 - \hat{\beta}_1 x_1 - \hat{\beta}_0)^2 + (y_2 - \hat{\beta}_1 x_2 - \hat{\beta}_0)^2 + \dots + (y_n - \hat{\beta}_1 x_n - \hat{\beta}_0)^2] =$$

$$-2(y_1 - \hat{\beta}_1 x_1 - \hat{\beta}_0) - 2(y_2 - \hat{\beta}_1 x_2 - \hat{\beta}_0) - \dots - 2(y_n - \hat{\beta}_1 x_n - \hat{\beta}_0) =$$

$$-2 \sum_{i=1}^n (y_i - \hat{\beta}_1 x_i - \hat{\beta}_0) = 0$$

$$\sum y_i - \hat{\beta}_1 \sum x_i - n\hat{\beta}_0 = 0$$

$$\hat{\beta}_0 = \frac{1}{n} \sum y_i - \frac{1}{n} \hat{\beta}_1 \sum x_i$$

Try to find β_1 by yourself

$$RSS = \sum_{i=1}^n (y_i - \bar{y} - \hat{\beta}_1(x_i - \bar{x}))^2$$

$$\frac{\partial}{\partial \hat{\beta}_1} RSS = \dots$$

Let's implement it

First of all, we need to generate data

```
def generate_data(b1, b0, size, x_range = (-10, 10), noise_mean = 0,
                  noise_std = 1):
    """
    input:
    b1, b0 - true parameters of data
    size - size of data, numbers of samples
    x_range - tuple of (min, max) x-values
    noise_mean - noise mean value
    noise_std - noise standard deviation

    output:
    data_x, data_y - data features
    """
    noise = np.random.normal(noise_mean, noise_std, size)
    rnd_vals = np.random.rand(size)
    data_x = x_range[1] * rnd_vals + x_range[0]*(1-rnd_vals)
    data_y = b1 * data_x + b0 + noise

    return data_x, data_y
```

Linear regression using LS

Secondly, we need function for data visualization

```
def animate(data_x, data_y, true_b1, true_b0, b1, b0, x_range = (-10,10),
            label="Least squares"):
    plt.scatter(data_x, data_y)
    plt.plot([x_range[0], x_range[1]],
             [x_range[0]*true_b1 + true_b0, x_range[1]*true_b1 + true_b0],
             c="r", linewidth=2, label="True")
    plt.plot([x_range[0], x_range[1]],
             [x_range[0]*b1 + b0, x_range[1]*b1 + b0],
             c="g", linewidth=2, label=label)
    plt.legend()
    plt.show()
```

Linear regression using LS

Now, let's implement prediction function

```
def least_squares(x, y):  
    """  
    input:  
    x, y - data features  
  
    output:  
    b1, b0 - predicted parameters of data  
    """  
    mean_x = x.mean()  
    mean_y = y.mean()  
  
    b1 = np.dot(y - mean_y, x - mean_x) / np.dot(x - mean_x, x - mean_x)  
    b0 = mean_y - b1*mean_x  
  
    return b1, b0
```


Linear regression using LS

Let's check how it works

```
### Parameters for data generation ###
true_b1 = 2.5
true_b0 = -7
size = 100
x_range = (0,10)
noise_mean = 0
noise_std = 1


# Generate the data
data_x, data_y = generate_data(true_b1, true_b0, size,
                               x_range = x_range,
                               noise_mean = noise_mean,
                               noise_std = noise_std)

# Predict data's parameters
b1, b0 = predict(data_x, data_y)

# Visualize the data
print("true b1 : {}\ntrue b0 : {}".format(true_b1, true_b0))
print("calculated b1 : {}\ncalculated b0 : {}".format(b1, b0))
animate(data_x, data_y, true_b1, true_b0, b1, b0, x_range=x_range)
```

Linear regression using LS

What we should get



Scikit-learn implementation



```
from sklearn.linear_model import LinearRegression

regression_model = LinearRegression()
# feed the linear regression with the train data to obtain a model.
regression_model.fit(X_train, y_train)
```

How to evaluate

R^2 metric

```
regression_model.score(X_test, y_test)
```

MSE

```
from sklearn.metrics import mean_squared_error  
import math
```

```
y_predict = regression_model.predict(X_test)  
regression_model_mse = mean_squared_error(y_predict, y_test)
```

How to predict

```
regression_model.predict([[4, 121, 110, 2800, 15.4, 81, 0, 1, 0]])
```

Homework explanation

On generated toy dataset

- Implement LR using GD and compare with LR based on LS
- Implement LR using Scikit-learn and compare with previous methods

On real dataset

- Implement multiple LR for the same using Scikit-learn

Let's take a look on a dataset

| mpg | cylinders | displacement | horsepower | weight | acceleration | model_year | origin | name |
|------|-----------|--------------|------------|--------|--------------|------------|--------|---------------------------|
| 18.0 | 8 | 307.0 | 130 | 3504 | 12.0 | 70 | 1 | chevrolet chevelle malibu |
| 15.0 | 8 | 350.0 | 165 | 3693 | 11.5 | 70 | 1 | buick skylark 320 |
| 18.0 | 8 | 318.0 | 150 | 3436 | 11.0 | 70 | 1 | plymouth satellite |
| 16.0 | 8 | 304.0 | 150 | 3433 | 12.0 | 70 | 1 | amc rebel sst |
| 17.0 | 8 | 302.0 | 140 | 3449 | 10.5 | 70 | 1 | ford torino |

We want to predict how many miles will car pass on one gallon of gasoline - MPG (miles per gallone)

What do you think about the dataset? What can we do with it ?

Does it have categorical predictors ? If yes, what we could do with them ?

Does any sample has all the fields (predictor values)?

Questions

- 1) Is there a relationship between mpg (response) and weight (predictor)? How we could check it?
- 2) What is the extent to which the model fits the data?
- 3) Is at least one of the predictors useful in predicting the response?
- 4) What could you say about dataset? What should we drop?

That's it for today! Questions?