

# Lab 8. Ensemble learning

Intro to Machine Learning  
Fall 2018, Innopolis University

# Lecture recap

- Ensemble learning
- Bagging
- Random Forests
- Boosting
- Adaboost

# Discussion

- How can we increase accuracy with ensemble learning?
- How can we reduce variance with ensemble learning?
- Discuss Bias-Variance Tradeoff.

Single classifier

$$a(x) = C(b(x))$$

$$b : X \rightarrow R$$

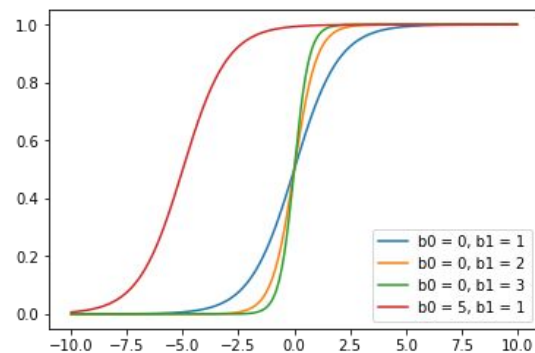
$$C : R \rightarrow Y$$

Single classifier

$$a(x) = C(b(x))$$

$$b : X \rightarrow R$$

$$C : R \rightarrow Y$$



$$b(x) = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$$

$$C(b(x)) = \begin{cases} 1 & b(x) > \text{threshold} \\ -1 & \text{otherwise} \end{cases}$$

## Ensemble Classifier

$$a(x) = C(F(b_1(x), \dots, b_T(x)))$$

$$F : R^T \rightarrow R$$

# Ensemble Classifier

$$a(x) = \textit{sign}(\alpha_1 b_1(x) + \dots + \alpha_T b_T(x))$$

- Train classifiers  $b_i(x)$
- Train weights  $\alpha_i$

# Bootstrap

From initial sample length  $\ell$  we are creating random sampling with replacement with the same length  $\ell$ .

Some objects are taken more, than once, some not included at all.

Part inside the new sample  $1 - e^{-1} \approx 0.632$  if  $\ell \rightarrow \infty$

```
a = range(10)
b = range(10, 20)

from sklearn.utils import resample
aa, bb = resample(a, b)

print(aa)
print(bb)
```

```
[9, 5, 1, 7, 3, 4, 7, 0, 3, 5]
[19, 15, 11, 17, 13, 14, 17, 10, 13, 15]
```



# Bagging

1. Generate a bootstrap sample.
2. Randomly select the subset of the predictors.
3. Train the basic algorithm
4. Repeat  $T$  times from 1 to 3

$$a(x) = \textit{sign} (b_1(x) + \dots + b_T(x))$$

Same weights for all classifiers

# Exercise

Recall from lecture:

Now you are asked to illustrate this on generated data:

- generate N (say, 10) sets of samples drawing them from Normal distribution (choose different *mean*, but the same *std* for each set)
- print variance for each set
- calculate the average of samples over all sets and print the resulting variance, compare

## Why Bagging Works?

- **Averaging** reduces variance
- Let  $Z_1, Z_2, \dots, Z_N$  be i.i.d random variables

$$\text{Var} \left( \frac{1}{N} \sum_i Z_i \right) = \frac{1}{N} \text{Var}(Z_i)$$

# Exercise. Hints

To sample from Normal distribution use:

```
samples = np.zeros((10, 100))  
for i in range(10):  
    samples[i, :] = np.array(np.random.normal(i, 0.5, 100))
```

At the end, you should get something similar to:

```
original variance for samples_0 0.25456984625140566  
original variance for samples_1 0.2776440860430631  
original variance for samples_2 0.27022972764984576  
original variance for samples_3 0.29694638366192955  
original variance for samples_4 0.23858918423833422  
original variance for samples_5 0.28246283016978035  
original variance for samples_6 0.23104593571693713  
original variance for samples_7 0.27474808156433594  
original variance for samples_8 0.20577265954609777  
original variance for samples_9 0.2784330299384325  
variance of the average 0.023107123035702837
```

# Benefits from bagging

- Outliers not included into some samples.
- Variance is reduced.

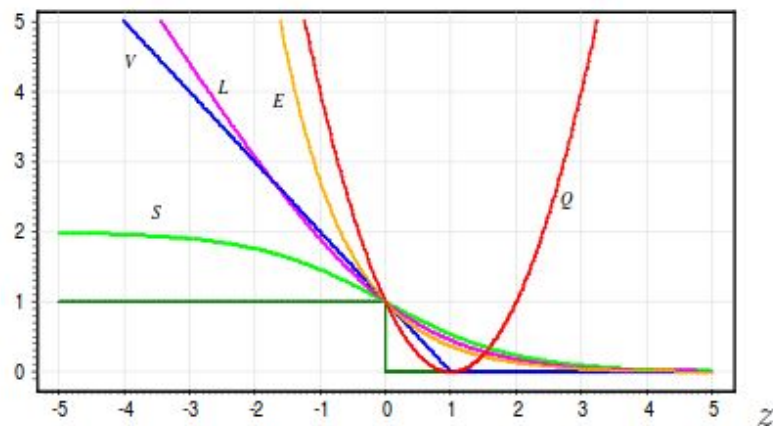
## Loss Function

$$L_T = \sum_{i=1}^l \left[ y_i \sum_{t=1}^T \alpha_t b_t(x_i) < 0 \right]$$

# Heuristics

- Greedy algorithm to optimize the function. Adding  $\alpha_t b_t(x)$  we assume  $\alpha_1 b_1(x), \dots, \alpha_{t-1} b_{t-1}(x)$  fixed. Optimizing only by  $\alpha_t b_t(x)$  parameters.
- Change loss function to differentiable one.

# Loss function



$$S(z) = 2(1 + e^z) - 1$$

*sigmoid*

$$L(z) = \log_2(1 + e^{-z})$$

*logarithmic*

$$V(z) = (1 - z)_+$$

*piecewise linear*

$$E(z) = e^{-z}$$

*exponential*

$$Q(z) = (1 - z)^2$$

*quadratic*

Adaboost

$$[y_i b(x_i) < 0] \leq e^{-y_i b(x_i)}$$



## Adaboost

$$\begin{aligned} L_T &\leq \tilde{L}_T = \sum_{i=1}^l \exp\left(-y_i \sum_{t=1}^T \alpha_t b_t(x_i)\right) = \\ &= \sum_{i=1}^l \underbrace{\exp\left(-y_i \sum_{t=1}^{T-1} \alpha_t b_t(x_i)\right)}_{w_i} e^{-y_T \alpha_T b_T(x_i)} \end{aligned}$$

# Algorithm

**Input:**  $X^\ell$ ,  $Y^\ell$  - training set.  $T$  - maximum number of base algorithms

**Output:** Base algorithms  $b_t(x)$  and their weights  $\alpha_t$

1. Initialize all weights  $w_i = 1/\ell$  for all  $i$  in  $1, \dots, \ell$ ;
2. For all  $t$  in  $1, \dots, T$ :

3.  $b_t = \arg \min_b N(b, W^l)$

$$\alpha_t = \frac{1}{2} \ln \frac{1 - N(b, W^l)}{N(b, W^l)}$$

$$N(b, W^l) = \sum_{i=1}^l w_i [b(x_i) \neq -y_i]$$

4. Recalculate weights  $w_i = w_i \exp(-\alpha_t y_i b_t(x_i))$

5. Normalize weights  $w_0 = \sum_{i=1}^l w_i$ ;  $w_i = w_i / w_0$ .

# sklearn

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier

algo = LogisticRegression()
model = AdaBoostClassifier(base_estimator=algo, n_estimators=10)

model.estimator_weights_
model.estimators_[j]
```

# Outliers

- High weights for misclassified objects.
- Adaboost tends to overfit to outliers.

# Outliers

- We can filter outliers using adaboost.
- Objects with high weights could be removed as outliers.
- The model should be retrained.
- Adaboost could be used only for outliers filtering. Another algorithm could be used after filtering.

# Discussion

What is the difference between **bagging** and **boosting**?

# Homework (start on the lab)

- In this homework you are not provided with any template.
- You will not be given precise instructions.
  - Understand yourself, how you will measure the quality, decide what is an outlier and so on.
- You can use any methods from **sklearn**.
- **Report** is required and very important in this HW. The quality will be graded.
- You should think yourself, what should and what shouldn't be in the report.
  - Describe your assumptions.
  - Analyze and explain results.
  - Justify your decisions.

# Homework

1. We are going to predict region based on country data. For simplicity let's make only two class classification - region is 'EUROPE' or not.
2. Download Countries dataset from Moodle.
3. Train adaboost classifier.
4. Find outliers in the dataset.
5. Retrain the model without outliers.



# References

<http://www.machinelearning.ru/wiki/images/0/0d/Voron-ML-Compositions.pdf>

Boosting the margin: a new explanation for the effectiveness of voting methods /  
R. E. Schapire, Y. Freund, W. S. Lee, P. Bartlett //

<https://web.stanford.edu/~hastie/Papers/samme.pdf>

[https://en.wikipedia.org/wiki/Bootstrapping\\_\(statistics\)](https://en.wikipedia.org/wiki/Bootstrapping_(statistics))